

BUSINESS PLAN

AI Liability Risk Assessment Agent

An Intelligent Underwriting Tool for the Emerging
AI Liability Insurance Market

Prepared: April 2026

Target: 2-Person Founding Team

Acquisition-Focused Strategy | 12–18 Month Horizon

CONFIDENTIAL

1. Executive Summary

THE OPPORTUNITY

Insurers are introducing broad AI exclusions across CGL, D&O, E&O, and cyber policies — creating a coverage vacuum. Specialty MGAs (Testudo, Armilla, Relm) are rushing to fill this gap, but lack the underwriting tooling to assess AI risk at scale. We build the risk scoring engine they desperately need.

We are building an AI agent that ingests an enterprise's AI governance documentation — model inventories, deployment architecture, bias audit reports, human-in-the-loop policies, incident logs — and produces a structured, evidence-based AI risk score for underwriters writing AI liability coverage.

Think of it as the cyber risk scoring platform (like SecurityScorecard or BitSight) but purpose-built for AI liability. Just as cyber underwriting evolved from paper questionnaires to automated continuous monitoring, AI liability underwriting will follow the exact same trajectory. We intend to be the platform that enables it.

Target exit: Acquisition by a data/analytics incumbent (Verisk, Guidewire, Moody's RMS) or a well-funded specialty MGA (Armilla, Testudo) within 12–18 months of launch, at \$5–25M based on comparable transactions.

2. Market Analysis

2.1 The AI Liability Insurance Market Is Exploding

AI liability insurance is a brand-new category being created in real time. The structural forces driving it are powerful and irreversible:

- **Coverage vacuum:** Verisk's ISO unit developed standard AI exclusion endorsements effective January 1, 2026, across multiple US states. WR Berkley and AIG have introduced or proposed exclusions barring any claim arising from AI use.
- **Demand surge:** A Geneva Association survey of 600 businesses found >90% would value insurance for AI risks, and >67% would pay at least 10% more for such coverage.
- **Litigation explosion:** AI-related lawsuits increased 81% YoY in Q1 2025, with average settlements around \$4M. The most common claims involve breach of contract, privacy invasion, and misrepresentation — not just IP/copyright.
- **Regulatory catalyst:** The EU AI Act, Colorado AI Act, and emerging state-level regulations are forcing companies to document AI governance — creating both insurable risk events and the very documentation our agent would ingest.

2.2 Who Is Writing AI Liability Policies Today

Company	Product	Capacity	Backing
Testudo	Gen AI liability policy	Up to \$10M	Lloyd's of London (Cohort 14)
Armillar AI	AI Liability + Performance Warranty	Up to \$25M	Lloyd's Coverholder; Chaucer, Swiss Re, Greenlight Re
Relm Insurance	NOVA AI, PONTA AI, RESCA AI	Varies by module	50-state MGA; specialty carrier
Munich Re	aiSure™ warranty	Performance guarantee	Own balance sheet (since 2018)
Vouch	AI-specific startup coverage	Varies	Series D, \$160M+ raised
CoverYourAI	Business interruption for AI ops	Emerging	Predictive underwriting models

These are our customers, not our competitors. They need underwriting infrastructure to assess risk at scale.

2.3 The Underwriting Tooling Gap

Today, AI liability underwriters assess risk via:

- Manual questionnaires about AI governance practices
- PDF-based reviews of applicant documentation
- External scanning of publicly disclosed AI incidents
- Red-teaming reports (when available, often not)

This is roughly where cyber insurance was in 2014 — before SecurityScorecard, BitSight, and similar platforms automated continuous risk monitoring. Testudo explicitly describes its approach as mining incident repositories and litigation data. Armilla reports conducting 500+ manual AI evaluations. Neither has a scalable, automated underwriting engine that ingests an applicant's governance posture in real time.

KEY INSIGHT

The AI liability insurance market is being born right now. The underwriters entering this space need tooling to scale beyond manual assessment. If we build the risk scoring infrastructure early, we become the data backbone of the entire category — the exact profile that attracts acquisition from data incumbents like Verisk (\$25.8B market cap, 76 acquisitions to date).

3. Product: What We Build

3.1 Core Agent Capabilities

Our product is a single intelligent agent that takes in an applicant's AI governance documentation and outputs a structured risk assessment for underwriters. The agent performs four sequential tasks:

Task 1 — Document Ingestion & Classification. The agent accepts uploads of AI governance documentation: model cards, model inventories, data processing impact assessments (DPIAs), incident response plans, bias audit reports, human oversight protocols, and deployment architecture diagrams. It classifies each document and identifies gaps (e.g., “no bias testing documentation provided”).

Task 2 — Governance Posture Extraction. From each document, the agent extracts structured data points across a governance taxonomy we define (mapped to EU AI Act, NIST AI RMF, ISO 42001, and Colorado AI Act requirements). Examples: number of high-risk AI systems deployed, presence of human-in-the-loop for consequential decisions, frequency of model retraining, documented incident response procedures.

Task 3 — External Risk Signal Enrichment. The agent cross-references the applicant against public data: AI-related litigation involving the company or its industry, regulatory enforcement actions, known AI incidents in the applicant's sector, and media sentiment around the company's AI practices.

Task 4 — Risk Score & Report Generation. The agent produces a structured risk report with: an overall AI risk score (1–100), sub-scores across governance dimensions, specific risk flags, coverage recommendations, and a confidence level for each assessment. Every finding links back to the source document and page.

3.2 Governance Taxonomy (Scoring Dimensions)

Dimension	Weight	What We Assess
Model Inventory & Classification	15%	Completeness of AI system registry; risk classification of each system (EU AI Act tiers)
Human Oversight	20%	HITL policies for consequential decisions; escalation paths; override documentation
Bias & Fairness Testing	15%	Frequency and methodology of bias audits; protected class impact analysis; remediation records
Data Governance	10%	Training data provenance; consent mechanisms; data retention; cross-border data flows

Incident Response	15%	Documented IR plans for AI failures; post-incident review process; notification procedures
Monitoring & Drift	10%	Model performance monitoring; drift detection; retraining triggers; A/B testing protocols
Regulatory Compliance	10%	Alignment with applicable frameworks (EU AI Act, NIST RMF, state laws); gap analysis
Third-Party AI Risk	5%	Vendor AI risk assessment; contractual protections; supply chain AI governance

3.3 Output for Underwriters

The underwriter receives a single-page summary score card plus a detailed 10–20 page report with:

- Overall risk score with industry benchmark comparison
- Red/amber/green flags per governance dimension
- Specific coverage recommendations (sublimits, exclusions, conditions)
- Document-linked evidence trail for every finding
- Suggested renewal monitoring triggers

4. Technical Architecture

4.1 System Overview

The system is intentionally simple — optimized for a two-person team to build and maintain:

Layer 1 — Ingestion Pipeline. PDF/DOCX parsing via a document processing service (e.g., Unstructured.io or AWS Textract). Documents are chunked, classified by type, and stored in a vector database (Pinecone or Weaviate) for retrieval.

Layer 2 — Governance Extraction Engine. An LLM-based (Claude API or GPT-4) extraction pipeline with domain-specific prompts for each governance dimension. Uses structured output schemas to extract typed data points from each document. Chain-of-thought reasoning ensures every extraction is traceable.

Layer 3 — External Risk Signals. A lightweight scraping/API layer that pulls AI litigation data (PACER, state court records), regulatory enforcement (EU AI Office, FTC, state AGs), and media sentiment. Testudo's public research on litigation volume is a proof point that this data is available and actionable.

Layer 4 — Scoring Model. A rules-based scoring engine (not ML — underwriters need deterministic, explainable outputs) that weights the extracted governance data against our taxonomy. The scoring model is configurable per carrier, allowing each underwriter to apply their own risk appetite.

Layer 5 — Report Generation. Templated report generation producing PDF output with the score card, detailed findings, and evidence citations. Designed to drop directly into an underwriting workbench or submission file.

4.2 Tech Stack (Lean)

Component	Choice	Rationale
Backend	Python (FastAPI)	Fast to build; strong PDF/NLP ecosystem
LLM	Claude API (Anthropic)	Best-in-class document reasoning; 200K context
Vector DB	Pinecone or Weaviate	Managed; scales without ops overhead
Document Processing	Unstructured.io	Handles messy insurance PDFs well
Frontend	React + Tailwind	Simple underwriter dashboard; file upload + results
Hosting	AWS (Lambda + S3 + RDS)	Serverless = low ops for 2 people
External Data	Custom scrapers + PACER API	AI litigation + regulatory data

4.3 Build Timeline

Phase	Timeline	Deliverable
Phase 0: Research & Design	Weeks 1–4	Governance taxonomy finalized; scoring model v1; 5+ underwriter interviews completed
Phase 1: Core Agent MVP	Weeks 5–12	Document ingestion + governance extraction for top 3 document types (model cards, bias audits, HITL policies)
Phase 2: Scoring & Reports	Weeks 13–18	Scoring engine + report generation; pilot with 2–3 specialty MGAs
Phase 3: External Signals	Weeks 19–24	Litigation data integration; regulatory monitoring; continuous risk signals
Phase 4: Scale & Acquire	Months 7–18	Expand document coverage; 5–10 MGA customers; position for acquisition

5. Go-to-Market Strategy

5.1 Customer Sequence

We sell to specialty MGAs and carriers writing AI liability — a universe of approximately 20–30 organizations today, growing rapidly. Our GTM is relationship-driven:

Step 1: Pilot with Testudo and/or Armilla. Both are early-stage, resource-constrained, and explicitly need better underwriting data. Testudo uses incident data but lacks governance posture assessment. Armilla does manual evaluations (500+ to date). Our agent automates and scales what they do by hand.

Step 2: Expand to Lloyd's syndicates exploring AI cover. Lloyd's Lab has been incubating AI liability products. Syndicates writing this coverage need underwriting tools. Our reports become a standard part of the submission package.

Step 3: Sell to incumbent carriers adding AI endorsements. As major carriers (AIG, Berkley, Zurich) develop their own AI liability products, they'll need risk assessment tooling. Our platform is carrier-configurable — each can apply their own risk weights.

5.2 Pricing Model

Per-assessment pricing (similar to how underwriters pay for Verisk reports or D&B credit checks):

- **Tier 1:** \$500–\$1,500 per assessment (basic governance score from uploaded docs)
- **Tier 2:** \$2,000–\$5,000 per assessment (full report with external signals + litigation data)
- **Enterprise:** \$50K–\$200K annual license for unlimited assessments (for MGAs processing volume)

At 100–300 assessments/year across 5–10 MGA customers, we reach \$500K–\$1.5M ARR within 12–18 months — sufficient for acquisition interest.

6. Acquirer Map

We have identified five tiers of potential acquirers, each with distinct strategic rationale:

6.1 Tier 1: Data & Analytics Incumbents (Most Likely)

Acquirer	Market Cap / Scale	Strategic Fit	Est. Acquisition Range
Verisk (VRSK)	\$25.8B market cap; \$3.07B revenue; 76 acquisitions	Already creating AI exclusion language via ISO; needs risk scoring tool to complement. Paid \$2.35B for AccuLynx, \$163M for SuranceBay in 2025 alone.	\$10–\$30M
Moody's / RMS	\$80B+ market cap	RMS does catastrophe modeling; AI risk is a new peril category. Moody's acquired RMS for \$2B; would bolt on AI risk scoring.	\$8–\$20M
Guidewire (GWRE)	\$15B+ market cap	Core system for P&C carriers; already invested in Sixfold. AI risk scoring would enrich their underwriting data layer.	\$5–\$15M

6.2 Tier 2: Specialty MGA Acquirers

Acquirer	Funding / Scale	Strategic Fit	Est. Acquisition Range
Armillia AI	\$25M raised; Lloyd's Coverholder	Currently doing manual evaluations at 500+ scale. Our agent automates their core underwriting workflow.	\$5–\$12M (acqui-hire + tech)
Testudo	Pre-seed (LocalGlobe, Goldman partners)	Needs governance scoring to complement their litigation data. Together = complete underwriting stack.	\$3–\$8M (acqui-hire + tech)
Relm Insurance	50-state MGA; 3 AI products	Broadest AI product suite but lacks automated applicant risk scoring.	\$5–\$10M

6.3 Tier 3: Consulting / Assurance Firms

The Big 4 (PwC, EY, Deloitte, KPMG) are building AI assurance practices — auditing AI systems for clients. Our risk scoring taxonomy and extraction engine could be repackaged as an AI assurance tool. Estimated range: \$5–\$20M.

6.4 Why They Buy (Not Build)

- Time-to-market: we're 6–12 months ahead with a working product and customer relationships

- Domain expertise: our governance taxonomy is built with underwriter input, not generic AI
- Data moat: every assessment we run enriches our benchmarking dataset — the more assessments, the more valuable our scores become
- Talent acquisition: for a 2-person team, the acqui-hire value alone justifies a \$3–\$5M floor

7. Competitive Positioning

Competitor	What They Do	What They Don't Do (Our Gap)
Testudo	Mines AI litigation data; scores risk from external signals	Does NOT ingest applicant's internal governance documentation
Armillar AI	Manual AI evaluations (500+); performance warranties	Does NOT have automated, scalable governance extraction
SecurityScorecard / BitSight	Cyber risk scoring from external signals	Zero AI-specific governance assessment capability
Big 4 AI Assurance	Manual AI audits for compliance	NOT integrated with insurance underwriting workflows; \$500K+ engagements
Generic LLM tools	Can read documents if prompted	No insurance-specific taxonomy; no scoring; no underwriter workflow integration

OUR MOAT

We sit at the intersection of AI governance expertise and insurance underwriting workflow — a niche no one else occupies. Testudo has the litigation data but not the governance assessment. Armilla has the insurance product but not the automated underwriting engine. SecurityScorecard has the scoring platform model but not the AI domain. We combine all three.

8. Financial Model

8.1 Revenue Projections

Metric	Month 6	Month 12	Month 18
MGA Customers	2–3 pilots	5–8 paying	8–12 paying
Assessments/Month	10–20	40–80	100–200

Avg Revenue/Assessment	\$1,000	\$1,500	\$2,000
Monthly Revenue	\$10K–\$20K	\$60K–\$120K	\$200K–\$400K
ARR (annualized)	\$120K–\$240K	\$720K–\$1.4M	\$2.4M–\$4.8M

8.2 Cost Structure (2-Person Team)

Item	Monthly Cost	Notes
Founder salaries (deferred/minimal)	\$5K–\$10K	Below market; equity-heavy
LLM API costs (Claude/GPT-4)	\$2K–\$8K	Scales with assessment volume
Infrastructure (AWS)	\$500–\$2K	Serverless keeps costs low
Document processing (Unstructured.io)	\$500–\$1K	Per-document pricing
Data sources (PACER, scrapers)	\$500–\$1K	Litigation monitoring
Total monthly burn	\$8.5K–\$22K	Can bootstrap or raise \$250K–\$500K pre-seed

9. Key Risks & Mitigations

Risk	Severity	Mitigation
Market timing: AI liability insurance doesn't scale as fast as expected	Medium	Even slow scaling creates a \$500M+ specialty market by 2028. Our per-assessment model means we only need 5–10 customers to be viable.
Incumbents build in-house (Verisk, Guidewire)	Medium	Building > buying takes 18+ months for incumbents. Our head start + customer relationships make acquisition cheaper than internal build. AccuLynx precedent: Verisk paid \$2.35B rather than build.
LLM accuracy on governance documents	High	We use structured extraction with human-in-the-loop review. Underwriters validate output. We start with high-confidence document types and expand.
Competitors pivot into our niche	Low	Testudo and Armilla are focused on their insurance products, not tooling. They're more likely to buy us than compete.
Regulatory uncertainty slows AI adoption	Low	Actually helps us — more regulation = more governance documentation = more demand for our scoring.

10. Immediate Next Steps (Weeks 1–4)

1. Conduct 5–10 discovery calls with specialty AI liability underwriters (Testudo, Armilla, Relm, Lloyd's syndicates). Validate the governance taxonomy and understand their current manual process.
2. Build a static prototype: take 3–5 sample AI governance documents, manually extract governance data, and produce a sample risk report. Use this as a sales tool in underwriter conversations.
3. Map the regulatory landscape: compile a definitive list of AI governance requirements across EU AI Act, NIST AI RMF, ISO 42001, Colorado AI Act, and NYC Local Law 144. This becomes the backbone of our scoring taxonomy.
4. Secure a design partner: get one MGA to commit to piloting the tool in exchange for early access and input on the scoring model. Testudo or Armilla are ideal — both are early-stage and resource-constrained.
5. Set up the technical foundation: document parsing pipeline, LLM extraction prompts, and scoring engine skeleton. Target: working demo on a single document type within 4 weeks.

End of Business Plan