

PHASE 0 — PRE-CODING PLAYBOOK

AI Liability Risk Assessment Agent

How AI Liability Underwriting Works Today
Test Data Strategy & Design Thinking Approach
Your First Step Before Writing a Single Line of Code

April 2026

CONFIDENTIAL

1. How AI Liability Underwriting Works Today

You won't have direct access to underwriters in the early days. This section documents exactly how the current manual process works, based on publicly available research from the key players. This is your substitute for customer interviews — it gives you the underwriter's perspective so you can build for their actual workflow.

1.1 The Current Manual Assessment Process

AI liability is a brand-new insurance class. Underwriters writing these policies come from cyber, tech E&O, or D&O backgrounds and are adapting their existing methods to a risk they're still learning to quantify. Here's what the process looks like today, reconstructed from Armilla AI, Testudo, Lockton Re, and industry reports:

Step 1 — Application & Questionnaire. The broker submits a standard application on behalf of their client. For AI liability, this includes a supplemental AI-specific questionnaire. These questionnaires are still being developed — most are adapted from cyber insurance questionnaires with AI-specific additions. Typical questions include: What AI systems does the applicant deploy? Are they using third-party foundation models or proprietary models? Is there human oversight for consequential decisions? Has the applicant conducted bias testing? What incident response procedures exist for AI failures?

Step 2 — Document Collection. The underwriter requests supporting documentation. For a mature applicant, this might include: an AI model inventory, model cards for high-risk systems, bias audit reports, data governance policies, human-in-the-loop protocols, incident response plans, and any regulatory compliance documentation (EU AI Act conformity assessments, NIST AI RMF alignment reports). For most applicants today, documentation is sparse or nonexistent — which itself is a major risk signal.

Step 3 — Manual Document Review. The underwriter (or a technical specialist) manually reads through submitted documents. They're looking for: completeness (are critical governance areas covered?), quality (are the policies specific and actionable, or vague boilerplate?), currency (when was the last bias audit? is the incident response plan current?), and red flags (gaps in documentation, inconsistencies between stated policies and actual practices).

Step 4 — External Research. The underwriter searches for public signals about the applicant's AI practices: active litigation involving AI (EEOC complaints, class actions, breach of contract claims), regulatory actions or investigations, media coverage of AI incidents, and the company's public AI ethics commitments versus actual behavior. Testudo specifically builds technology to mine AI litigation data and incident repositories for this purpose.

Step 5 — Risk Assessment & Pricing. The underwriter synthesizes everything into a risk judgment. Unlike traditional lines where actuarial data drives pricing, AI liability has almost no historical loss data. Underwriters are pricing based on governance posture, industry exposure, and analogies to cyber risk. Armilla’s approach evaluates “specific product dimensions and usage contexts” and develops proprietary datasets to quantify risk. The Lockton Re/Armilla report specifically recommends assessing a “target model metric” where each AI model is underwritten on its individual merits.

Step 6 — Decision & Terms. The underwriter decides: accept (with specific terms, sublimits, and conditions), decline, or request additional information. For accepted risks, terms might include: required human oversight for specific decision types, mandatory bias testing frequency, incident notification requirements, and documentation maintenance obligations.

KEY INSIGHT: THE BOTTLENECK IS STEPS 3 AND 4

Steps 1–2 are standard insurance workflow. Steps 5–6 require human judgment and will always involve an underwriter. But Steps 3 (document review) and 4 (external research) are where underwriters spend the most time and where accuracy varies most between reviewers. This is exactly where your agent sits — automating the document analysis and external signal gathering, then presenting structured findings for the underwriter’s judgment.

1.2 What Underwriters Actually Evaluate (The Scoring Dimensions)

Based on the Lockton Re/Armilla report, Armilla’s assessment methodology, the Trustible/Armilla partnership framework, and NIST AI RMF, here are the specific dimensions underwriters care about:

Dimension	What Underwriters Look For	Source
AI System Inventory	How many AI systems? What risk tier? Where deployed? What decisions do they make?	Lockton Re/Armilla report; NIST GOVERN 1.1
Model-Level Risk	Individual model metrics: accuracy, robustness, explainability, fairness per model	Armilla Verified assessment; Lockton Re "target model metric" approach
Human Oversight	HITL for consequential decisions; escalation paths; override authority; review cadence	EU AI Act Art. 14; NIST GOVERN 1.3; Armilla questionnaire
Bias & Fairness	Testing methodology; protected classes; remediation; testing frequency; NYC LL144 compliance	Armilla bias detection; NIST MEASURE 2.11; NYC LL144
Data Governance	Training data provenance; consent; retention; cross-border flows; privacy compliance	NIST MAP 2; GDPR/DPIA requirements

Incident Response	AI-specific IR plan; notification procedures; post-incident review; SLAs	NIST MANAGE 4.1; Airmic guidance
Monitoring & Drift	Performance monitoring; drift detection; retraining triggers; continuous evaluation	Armilla ongoing monitoring; NIST MEASURE 3
Regulatory Alignment	EU AI Act conformity; NIST RMF alignment; state law compliance (CO, NYC, IL)	NIST GOVERN 1.1; EU AI Act risk classification
Third-Party AI Risk	Vendor model dependencies; shared foundation model risk; supply chain governance	Lockton Re systemic risk analysis; NIST MAP 5
Governance Maturity	Board oversight; AI ethics committee; training programs; culture of responsible AI	NIST GOVERN 1; Armilla governance signals via Trustible

1.3 How Current Players Do Assessments (Without Talking to Them)

Here's what you can learn from each player's public materials about their methodology:

Armilla AI — The Most Detailed Public Methodology

- **Assessment approach:** Third-party AI verification and red-teaming. They run automated test suites against models for accuracy, robustness, explainability, and fairness. Results feed an interactive dashboard.
- **Underwriting input:** Their assessments evaluate “training data, testing methods, and usage scenarios.” Strong governance signals from tools like Trustible “inform Armilla underwriting, enabling better terms as risk posture improves — similar to safe-driver effects in auto insurance.”
- **Coverage scope:** AI model errors, hallucinations, regulatory violations, data leakage, bias/discrimination, agent failures, IP infringement. Up to \$25M per company.
- **What they DON'T have:** Automated ingestion of governance documentation. Their 500+ evaluations are hands-on assessments, not document parsing. This is your gap.

Testudo — Litigation Data Focus

- **Assessment approach:** External-facing. They compile and analyze real-time AI litigation data, regulatory developments, and incidents. Their platform mines incident repositories to assess risk severity.
- **Key data:** AI lawsuits increased 81% YoY in Q1 2025. Top states: CA, IL, NY, TX, DE. Most common claims: breach of contract, invasion of privacy, misrepresentation. Average settlements ~\$4M.

- **Underwriting method:** Does NOT require deep integration with a company's AI systems. Uses external signals only, which makes it simpler to adopt but less comprehensive than governance document analysis.
- **What they DON'T have:** Internal governance posture assessment. They see what's public (lawsuits, incidents) but not what's inside the company (policies, testing, oversight). This is your gap.

Lockton Re / Armilla Joint Report — Industry Blueprint

- **Key recommendation:** Each AI model should be “underwritten on its individual merits, based on factors such as the industry, context of the outputs, any underlying foundation model, the version deployed and the use case.”
- **Systemic risk warning:** Shared AI infrastructure (everyone using the same foundation models) creates correlated risk. A single model failure could trigger claims across thousands of policyholders simultaneously.
- **What they call for:** Underwriters must evaluate “portfolio-level exposure concentration through shared model dependencies, architectural vulnerability to coordinated attacks, and the capacity to detect failures before substantial liability accrues.”

Trustible + Armilla Partnership — The Governance-to-Insurance Pipeline

- **How it works:** Trustible provides the governance platform (inventories, controls, monitoring). Armilla uses those governance signals to underwrite coverage. Better governance = better insurance terms.
- **Assessment process:** “A joint discovery identifies your AI systems, applicable obligations, and loss scenarios. Trustible sets up inventories, controls, and monitoring; Armilla evaluates governance/reliability signals and structures coverage through an insurance broker.”
- **Your positioning:** You replace the “joint discovery” step with an automated agent that can ingest any company's governance docs and produce the same structured output — without requiring the company to be a Trustible customer.

YOUR COMPETITIVE POSITIONING IN ONE SENTENCE

Testudo sees outside the company (litigation, incidents). Armilla tests inside the company (model performance). You read the company's own documentation and extract what the underwriter needs. Combined, all three layers would give an underwriter a complete picture — but your layer is the one that doesn't exist yet as an automated tool.

2. Test Data Strategy (Design Thinking Approach)

This follows the five-phase design thinking framework: Empathize, Define, Ideate, Prototype, Test. Since you won't have underwriter access initially, we substitute publicly available research for empathy interviews and use synthetic-but-realistic test data.

PHASE 0A — Week 1, Days 1–3

Empathize — Build Underwriter Mental Model

Since you can't interview underwriters directly, build your understanding from these public sources:

Required Reading List (Do This Before Anything Else)

Document	Source	What You'll Learn
Lockton Re / Armilla AI: AI Risk Classification Report	carriermanagement.com (Feb 2026)	How underwriters should classify AI risk; the "target model metric" approach; systemic risk concerns
Armilla: Quasi-Regulatory Landscape Report	armilla.ai/resources	How AI liability policies actually work; what insurers assess; gaps between coverage and risk
NIST AI RMF 1.0 (AI 100-1)	nvlpubs.nist.gov	The governance framework your scoring taxonomy maps to; GOVERN/MAP/MEASURE/MANAGE functions
NIST AI 600-1: Generative AI Profile	nvlpubs.nist.gov	GenAI-specific risk factors and suggested management actions
Airmic: AI — Will Your Insurance Payout?	airmic.com	UK risk manager perspective on AI coverage gaps; practical governance checklist
Verisk ISO AI Exclusion Language	Search "Verisk AI exclusion endorsement 2026"	The exact exclusion language driving the coverage vacuum your product fills
Testudo website + Lloyd's Lab materials	testudo.co / lloyds.com/lloyds-lab	How litigation-driven underwriting works; their data approach
NYC Local Law 144 Bias Audits	Search "NYC LL144 bias audit published 2024 2025"	Real-world bias audit reports you can use as test documents

Industry Conference Recordings (Free Underwriter Perspectives)

- ITC Vegas 2025 recordings — multiple sessions on AI liability insurance
- Lloyd's Lab Demo Day recordings — Testudo and Armilla both presented
- Podcast: Insurance Thought Leadership — episodes on AI risk with carrier executives
- Podcast: Insurtech Insider — interviews with AI liability MGA founders
- YouTube: search “AI liability insurance underwriting 2025 2026” for panel discussions

PHASE 0B — Week 1, Days 3–5

Define — Assemble Real Test Documents

Collect real, publicly available AI governance documents into your test corpus. These are not synthetic — they're actual production documents from real companies:

Source 1: Public Model Cards (Real Production Docs)

Company	Document	Where to Find It	Document Type
Anthropic	Claude model card + system prompts	anthropic.com/research + docs.anthropic.com	model_card
OpenAI	GPT-4 system card + Model Spec	openai.com/research	model_card, regulatory_compliance
Google	Gemini technical reports	deepmind.google/research	model_card
Meta	Llama model cards	llama.meta.com	model_card
xAI	Grok 4 model card (46-page PDF)	data.x.ai	model_card
HuggingFace	Thousands of community model cards	huggingface.co/models	model_card

Source 2: Published Bias Audits (Real Compliance Docs)

NYC Local Law 144 requires companies using automated employment decision tools (AEDTs) to publish annual bias audits. These are real, legally required documents:

- Search: “bias audit report AEDT 2024 2025 filetype:pdf”
- Companies with published audits: HireVue, Harver (formerly Pymetrics), iCIMS, Paradox, Greenhouse, Workday
- These are perfect test documents for the bias_fairness extraction dimension
- Each audit contains: methodology, protected class analysis, impact ratios, and recommendations

Source 3: Governance Frameworks & Templates (Real Policy Docs)

Document	Source	Document Type
NIST AI RMF Playbook (125 pages)	airc.nist.gov/airmf-resources/playbook	regulatory_compliance
Singapore Model AI Governance Framework	pdpc.gov.sg	regulatory_compliance
WEF Model AI Governance Framework	weforum.org	regulatory_compliance
UK ICO AI Guidance + DPIA templates	ico.org.uk	dpia, data_governance
EU AI Act compliance templates	artificialintelligenceact.eu	regulatory_compliance
IEEE 7000 series standards (summaries)	standards.ieee.org	regulatory_compliance
ISO 42001 overview documents	iso.org	regulatory_compliance
Google Responsible AI Practices	ai.google/responsibility	hitl_policy, bias_audit
Microsoft Responsible AI Standard (v2)	microsoft.com/ai/responsible-ai	hitl_policy, incident_response

Source 4: AI Incident Data (Real External Signals)

- **AI Incident Database (incidentdatabase.ai):** 600+ documented AI failures with descriptions, dates, involved parties, and severity. This is your external_signals test data.
- **AIAAIC Repository (aiaaic.org):** Curated database of AI incidents, controversies, and harms with source links.
- **Testudo's public research:** They've published that AI lawsuits increased 81% YoY, top states are CA/IL/NY/TX/DE, and average settlements are ~\$4M. Use these benchmarks for your scoring calibration.

PHASE 0C — Week 2, Days 1–3

Ideate — Create Three Fictional Company Personas

Create three fictional companies at different risk levels. For each persona, mix real public documents with Claude-generated synthetic documents to create a complete submission package. This gives you realistic test data AND a compelling demo.

Persona 1: “GreenScore AI” — Low Risk (Target Score: 78–90)

Attribute	Detail
-----------	--------

Industry	Fintech — ESG scoring and portfolio analytics
AI Systems	3 models: ESG text classifier, portfolio optimizer, risk predictor
Employees	200; dedicated AI governance team of 4
Governance Maturity	High — NIST AI RMF aligned; board AI ethics committee; quarterly bias audits
Key Risk Factors	Uses third-party foundation model (GPT-4) for text classification; EU clients trigger AI Act obligations
Litigation History	None
Documents to Create	Model inventory (complete), 3 model cards, bias audit report (recent), HITL policy, incident response plan, DPIA, data governance policy, NIST RMF self-assessment

Claude prompt to generate docs: “Generate a realistic AI model inventory document for GreenScore AI, a fintech company with 200 employees that uses 3 AI models for ESG scoring. The inventory should be thorough and well-maintained, following NIST AI RMF guidelines. Include: model name, purpose, risk classification, deployment date, last review date, data sources, human oversight level, and monitoring status for each model. Make it look like a real internal document, not a template.”

Persona 2: “QuickHire Inc” — Medium Risk (Target Score: 40–60)

Attribute	Detail
Industry	HR Tech — AI-powered resume screening and candidate matching
AI Systems	2 models: resume screener, candidate-job matcher; plus ChatGPT integration for interview scheduling
Employees	85; no dedicated AI governance role
Governance Maturity	Partial — has a model card but it’s 18 months outdated; bias audit done once (initial launch); HITL only for “senior roles”; no formal incident response for AI failures
Key Risk Factors	NYC LL144 applies; one pending EEOC complaint; bias audit is stale; ChatGPT integration has no governance wrapper
Litigation History	1 active EEOC complaint alleging disparate impact in resume screening
Documents to Create	Outdated model card, initial bias audit (18 months old), partial HITL policy, basic data retention policy, NO incident response plan, NO formal model inventory

Claude prompt: “Generate a model card for QuickHire Inc’s resume screening AI. Make it look realistic but with clear governance gaps: the ‘last updated’ date should be October 2024, the bias testing section should say ‘initial validation completed at launch’ with no follow-up, the

limitations section should be vague, and there should be no mention of NYC Local Law 144 compliance. It should look like a document written by engineers without governance guidance.”

Persona 3: “AutoClaim Pro” — High Risk (Target Score: 12–30)

Attribute	Detail
Industry	Insurtech — AI-powered claims auto-adjudication
AI Systems	5+ models: claims triage, fraud detection, damage estimation, auto-denial engine, customer chatbot
Employees	150; CTO “owns AI” with no formal governance structure
Governance Maturity	Minimal — no model inventory, no bias testing, no formal HITL for claims under \$10K, no incident response plan, no DPIA despite processing personal health data
Key Risk Factors	Auto-denies claims under \$10K without human review; two active lawsuits (wrongful denial class action + individual bad faith claim); Colorado AG inquiry about AI decision-making in claims; uses 3 different foundation models with no vendor risk assessment
Litigation History	2 active lawsuits, 1 regulatory inquiry
Documents to Create	Brief product overview (marketing-style, not governance), basic privacy policy (generic, not AI-specific), one-page “AI ethics statement” (vague commitments), NO model cards, NO bias audits, NO incident response

Claude prompt: “Generate a one-page AI ethics statement for AutoClaim Pro, an insurtech company. Make it a corporate PR document with vague commitments like ‘we are committed to responsible AI’ and ‘we believe in fairness and transparency’ but with NO specific policies, NO testing methodology, NO governance structure, and NO mention of human oversight for claims decisions. It should read like a marketing page, not a governance document.”

WHY THREE PERSONAS MATTER

When you demo for potential acquirers or MGA customers, you run all three through the system live. GreenScore shows the tool working well on a mature company. QuickHire shows it catching real governance gaps. AutoClaim shows it flagging severe risk. The contrast between the three scores is more convincing than any slide deck.

PHASE 0D — Week 2, Days 3–5

Prototype — Build the Test Corpus

Create the actual folder structure and documents:

```

test_data/
├── real_documents/                                # Actual public documents
│   ├── model_cards/
│   │   ├── anthropic_claude_model_card.pdf
│   │   ├── openai_gpt4_system_card.pdf
│   │   └── xai_grok4_model_card.pdf
│   ├── bias_audits/
│   │   ├── hirevue_ll144_audit_2024.pdf
│   │   └── harver_ll144_audit_2024.pdf
│   ├── frameworks/
│   │   ├── nist_ai_rmf_100-1.pdf
│   │   ├── nist_ai_600-1_genai_profile.pdf
│   │   └── nist_ai_rmf_playbook.pdf
│   └── incidents/
│       └── ai_incident_database_export.csv
├── company_a_greenscore/                          # Low risk persona
│   ├── model_inventory.pdf
│   ├── model_card_esg_classifier.pdf
│   ├── model_card_portfolio_optimizer.pdf
│   ├── model_card_risk_predictor.pdf
│   ├── bias_audit_q1_2026.pdf
│   ├── hitl_policy.pdf
│   ├── incident_response_plan.pdf
│   ├── dpia_esg_scoring.pdf
│   ├── data_governance_policy.pdf
│   └── nist_rmf_self_assessment.pdf
├── company_b_quickhire/                          # Medium risk persona
│   ├── model_card_resume_screener.pdf            # outdated
│   ├── initial_bias_audit_2024.pdf               # stale
│   ├── hitl_policy_senior_roles.pdf              # partial
│   ├── data_retention_policy.pdf                 # basic
│   └── # NOTE: no incident response, no model inventory
├── company_c_autoclaim/                          # High risk persona
│   ├── product_overview.pdf                      # marketing doc
│   ├── privacy_policy.pdf                        # generic
│   ├── ai_ethics_statement.pdf                  # vague
│   └── # NOTE: almost no real governance docs
└── expected_outputs/                             # Ground truth for testing
    ├── company_a_expected_score.json
    ├── company_b_expected_score.json
    └── company_c_expected_score.json

```

Ground Truth Files

For each persona, create a JSON file with the expected extraction results and scores. This becomes your automated test suite:

```
// company_b_expected_score.json
{
  "organization": "QuickHire Inc",
  "expected_overall_score": { "min": 38, "max": 62 },
  "expected_risk_tier": "medium",
  "expected_dimension_scores": {
    "model_inventory": { "min": 15, "max": 30,
      "key_flags": ["no_formal_inventory"] },
    "human_oversight": { "min": 35, "max": 55,
      "key_flags": ["partial_hitl_coverage", "senior_roles_only"] },
    "bias_fairness": { "min": 20, "max": 40,
      "key_flags": ["stale_audit", "no_ongoing_testing"] },
    "incident_response": { "min": 0, "max": 10,
      "key_flags": ["no_ir_plan"] },
    "regulatory_compliance": { "min": 25, "max": 45,
      "key_flags": ["ll144_noncompliance_risk"] }
  },
  "expected_critical_flags": [
    "Bias audit is 18+ months old",
    "No incident response plan for AI failures",
    "No human oversight for non-senior role decisions",
    "Active EEOC complaint not addressed in documentation",
    "ChatGPT integration has no governance wrapper"
  ]
}
```

PHASE 0E — Week 2, Day 5 + Ongoing

Test & Validate — The Golden Flow

Your golden flow is the end-to-end demo that proves the system works. It's also your primary automated integration test. Here's the exact flow:

1. **Upload QuickHire's 4 documents** to the system via the API or UI
2. **System classifies each document:** model_card (confidence: 0.92), bias_audit (confidence: 0.95), hitl_policy (confidence: 0.88), data_governance (confidence: 0.85)
1. **System extracts governance data per dimension** — for dimensions with no matching documents (incident_response, model_inventory), the system flags “no documentation provided” as a critical finding
1. **System gathers external signals:** 1 active EEOC complaint, NYC LL144 regulatory risk, 3 media articles about AI hiring bias in the HR tech sector
1. **Scoring engine produces dimension scores** and overall score (~50, medium risk). Red flags: stale bias audit, missing IR plan, partial HITL. Amber flags: outdated model card, no formal inventory. Green flags: data retention policy exists.

1. **PDF report generates** with radar chart, dimension breakdowns, evidence citations, and recommended terms (e.g., “require updated bias audit within 60 days as condition of coverage”)

Automated validation: After each run, compare actual output against the expected_outputs JSON. Scores should fall within expected ranges. All critical flags should be detected. Any extraction that falls outside expected bounds triggers an alert for manual review.

3. Week-by-Week Phase 0 Execution Plan

Day	Task	Output
Day 1	Read the Lockton Re/Armilla report + NIST AI RMF 1.0. Take notes on what underwriters actually evaluate.	Annotated notes with scoring dimension candidates
Day 2	Read Armilla’s public methodology + Testudo’s website + Airmic guidance. Map to your taxonomy.	Draft governance taxonomy (8–10 dimensions with weights)
Day 3	Download real model cards (Anthropic, OpenAI, xAI, Meta). Download 2–3 NYC LL144 bias audits.	real_documents/ folder populated with 8–10 real PDFs
Day 4	Use Claude to generate GreenScore AI’s full document set (10 docs). Review and refine for realism.	company_a_greenscore/ folder with 10 synthetic docs
Day 5	Generate QuickHire Inc docs (4 docs with deliberate gaps) and AutoClaim Pro docs (3 weak docs).	company_b_quickhire/ and company_c_autoclaim/ populated
Day 6	Create expected_outputs JSON files for all 3 personas. Define score ranges and critical flags.	expected_outputs/ folder with ground truth for automated testing
Day 7	Manually walk through the QuickHire assessment as if you were an underwriter. Document every judgment call.	Written golden flow walkthrough — becomes your agent’s logic spec
Day 8	Set up the Git repo, write CLAUDE.md, create the folder structure, docker-compose for local Postgres.	Project skeleton ready. You can start Phase 1 coding on Day 9.

THE MANUAL WALKTHROUGH IS THE MOST IMPORTANT STEP

On Day 7, sit down with QuickHire’s docs and pretend you’re an underwriter. Read each document. Write down every observation. Score each dimension by hand. Note where you’re confident and where you’re uncertain. This manual walkthrough becomes the specification for your extraction prompts and scoring rules. If you can’t score it by hand, your agent can’t score it either.

4. Getting Underwriter Feedback Without Direct Access

You said you won't have access to these people. Here's how to validate without formal interviews:

4.1 LinkedIn Research & Cold Outreach

- **Search:** "AI liability underwriter" OR "AI risk underwriter" OR "tech E&O underwriter" on LinkedIn. The universe is ~200–500 people globally. Many will have posts about AI risk.
- **Engage with content:** Comment thoughtfully on posts by underwriters at Testudo, Armilla, Relm, Lloyd's syndicates. Build visibility before asking for anything.
- **Cold DM template:** "Hi [name], I'm building a tool to help AI liability underwriters assess governance posture from submitted documentation. Would you have 15 minutes to look at a sample risk report and tell me if it's missing anything? Happy to share the output." The offer of free output makes this a value exchange, not a cold ask.

4.2 Insurance Industry Forums & Events

- **ITC Vegas / ITC London:** Insurtech conference. Testudo and Armilla exhibit here. Attend the AI liability sessions and network with underwriters directly.
- **Lloyd's Lab Demo Days:** Open to the public. Attend and meet the people building AI liability products.
- **r/InsuranceProfessionals + r/actuary:** Reddit communities where underwriters discuss emerging risks. Post your sample report and ask for feedback.
- **Insurance industry Slack communities:** Insurtech Slack, The Binder (specialty insurance), and similar groups.

4.3 The "Sample Report" Strategy

Your best validation tool is a finished sample report. Before writing any code:

3. Run the QuickHire manual assessment by hand
4. Design the report format in Figma or Google Docs
5. Include: score card, dimension breakdown, flags, evidence trail, coverage recommendations

6. Post it in insurance communities with: “We’re building an automated AI governance scoring tool for underwriters. Here’s what a sample report looks like. What’s missing? What would change your pricing decision?”
7. Share it with Armilla and Testudo contacts. Even if they don’t respond, they’ll see it and remember you when you have a working product

4.4 Proxy Validation: Use AI Governance Consultants

If you can’t reach underwriters, AI governance consultants are the next best thing. They advise companies on the exact documentation your tool will ingest:

- Trustible (trustible.ai) — AI governance platform; they know what documentation companies actually produce
- Holistic AI — AI risk management; they conduct audits that produce the documents you’ll ingest
- Credo AI — AI governance platform; they define what “good” governance looks like
- Big 4 AI assurance teams — PwC, EY, Deloitte all have public AI governance frameworks you can benchmark against

These consultants sit between the companies producing governance docs and the underwriters consuming them. They can validate both sides of your product: “Is our extraction accurate?” (they know what the docs should say) and “Is our scoring reasonable?” (they know what underwriters care about).

*This document is Phase 0 of the AILRA Technical Architecture.
Complete this before opening Claude Code. Then proceed to Phase 1 in the Technical Architecture document.*